

Project Report

Reimagining Artificial Intelligence: The Credibility Dashboard

Introduction and Project Context

This project was undertaken as part of the MGMT 522 New Product Development course with the objective of identifying a real, unmet consumer problem related to artificial intelligence usage and developing a feasible, differentiated product solution. The project focused on reimagining how users interact with generative AI systems, particularly in high-stakes academic and professional contexts where accuracy, trust, and credibility are critical. The team approached the project using a structured new product development framework, combining consumer research, ideation, validation, design, and feasibility analysis to arrive at a final recommendation.

Problem Identification: The Credibility Crisis in AI

The project began with identifying a growing credibility crisis among students who use generative AI tools for academic learning, research, and writing. While students increasingly rely on AI to support their academic work, they struggle to trust the outputs due to a lack of transparency, unclear sourcing, and uncertainty around accuracy. The team recognized that the core issue was not simply the presence of incorrect answers, but rather the inability of AI systems to communicate how confident they are in their responses and where the information originates from. This lack of visibility creates hesitation, anxiety, and repeated manual fact-checking, ultimately reducing the perceived value of AI tools.

The problem was reframed as an emotional trust crisis rather than a purely technical accuracy problem. Students reported fear of submitting AI-assisted work, embarrassment when incorrect information surfaced, and a constant need to verify answers independently. This reframing allowed the team to move beyond improving AI accuracy and instead focus on restoring user confidence and peace of mind.

Consumer Research and Insight Development

To validate the problem, the team conducted primary qualitative research through student interviews. Interview responses revealed consistent themes of anxiety, mistrust, and frustration. Students reported using AI as a starting point rather than a final authority, expressing discomfort with submitting work generated by AI without verification. A recurring concern was the absence

of visible sources, which forced users to manually cross-check information through external searches.

These insights were supported by secondary research from academic journals and industry publications. Prior studies highlighted transparency as a key trust-maker in AI systems and emphasized that users are less likely to trust AI outputs when they cannot understand how conclusions are reached. Industry data further confirmed that a significant majority of users double-check AI responses and lose trust when sources are not visible. Together, these findings provided strong evidence that credibility and transparency represent a substantial unmet need in current AI products.

Defining User Values and Needs

Based on the research findings, the team identified a set of core cognitive and emotional values that the solution needed to address. These included clarity, risk reduction, quality, truthfulness, emotional security, intellectual integrity, and peace of mind. This step was critical in translating abstract feelings of anxiety and mistrust into actionable product principles. By grounding the solution in these values, the team ensured that subsequent design and feature decisions remained closely aligned with genuine user needs rather than technical novelty.

Ideation and Concept Evaluation

During the ideation phase, the team explored multiple solution pathways before narrowing down to two primary concepts: the Credibility Dashboard and the Credibility Lens. The Credibility Lens focused on giving users control over what types of sources they considered credible, while the Credibility Dashboard emphasized providing transparent credibility indicators directly alongside AI outputs.

A structured comparison of the two concepts revealed that the Credibility Dashboard better addressed the broad trust gap across AI users. It required less effort from users, reduced cognitive load, and delivered reassurance rather than forcing users to actively define credibility standards. Additionally, the dashboard concept demonstrated higher feasibility by building on existing AI sourcing and confidence estimation models. Based on these considerations, the team selected the Credibility Dashboard as the final concept.

Product Concept: The Credibility Dashboard

The Credibility Dashboard was designed as an integrated transparency layer within ChatGPT that allows users to assess the trustworthiness of AI responses at a glance. The product

enables users to understand not only what the AI says, but why it says it. Key features include a confidence score that reflects the model's certainty, a visual breakdown of sources by type, an expert consensus indicator to show alignment across trusted sources, and a timeline view to assess how current the information is.

Each feature was directly mapped to a specific user concern identified during research. Anxiety about accuracy was addressed through confidence scoring, uncertainty about origins was addressed through source visualization, and concerns about outdated information were addressed through expert consensus tracking. This clear mapping ensured strong problem–solution fit.

Positioning and Market Differentiation

The Credibility Dashboard was positioned for students and professionals who rely on AI but dislike playing the role of fact-checker. Unlike existing AI tools such as ChatGPT or Gemini, which provide confident but opaque answers, the Credibility Dashboard differentiates itself by making credibility visible and interpretable. A perceptual map further illustrated that the product occupies a unique space characterized by high transparency and deep credibility insight, an area largely unaddressed by current competitors.

Design and User Experience Development

From a design perspective, the team deliberately chose to integrate the dashboard within the familiar ChatGPT interface rather than creating a standalone product. This decision reduced friction, improved adoption likelihood, and leveraged existing user trust in the interface. Design elements such as circular confidence gauges, source breakdown charts, expandable source cards, and emoji-based feedback sliders were selected to make complex credibility information easily scannable and emotionally approachable. Each design decision was supported by a clear rationale grounded in usability and behavioral psychology.

Technical Feasibility and Implementation

The team proposed a realistic and scalable technical architecture for the Credibility Dashboard. The system was designed as a wrapper around an existing large language model rather than a new model. A frontend built with modern web technologies would display the dashboard, while a middleware layer would coordinate between the AI model and a credibility analyzer service. The backend credibility analyzer would evaluate AI responses using semantic search against trusted sources and assign reliability scores. User feedback would be stored and used to continuously

refine confidence calculations. This architecture demonstrated that the product could be implemented using existing technologies with manageable complexity.

Validation, Testing, and Results

To validate the proposed solution, the team analyzed external market data and conducted user testing. Industry statistics confirmed widespread credibility concerns, with a majority of users reporting that they double-check AI outputs and lose trust when sources are hidden. User testing results showed a strong preference for the Credibility Dashboard over the Credibility Lens. Participants believed the dashboard would increase trust, save time, reduce fact-checking behavior, and be especially valuable for academic research. These results provided strong evidence of market desirability.

Rollout Strategy and Final Recommendation

The team proposed a phased rollout strategy, beginning with a prototype to demonstrate proof of concept, followed by a closed beta to gather feedback, a public launch through integrations or extensions, and ongoing refinement using user feedback. This roadmap reflected an understanding that credibility is dynamic and must continuously evolve.

Based on the strength of the research, validation results, feasibility analysis, and user response, the final recommendation was to launch the Credibility Dashboard as the primary solution. The project successfully demonstrated a complete new product development cycle, from identifying a deeply human problem to delivering a practical, differentiated, and validated solution.